Tool Giving Intuitive and Analytical Index Suggestions

By:

Jayant Shelke

Mustafa Badshah

Muthaiah Ramanathan

Project Specifications

- Tool that analyses a workload of queries that are executed against a Database Management System and suggests indexes to improve performance.

- Console based application developed in Visual Studio, Dot Net Framework 4.5
- Back-end development database – Microsoft SQL Server
- Language Used – C#
- Sample Database – AdventureWorks2016
- Collaboration Base - Github

Motivation

- Improving the performance of database transactions remains a challenge in practice.
- Challenge is associated with how data records are organized in a storage device and how quickly those records can be accessed.
- Common task undertaken by DB Administrators is selection of indexes to speed up performances of queries submitted to DBMS.
- Difficult to manually tune performance for extremely large database schema as possible configurations grow exponentially.

Challenges

- Index Selection is not straightforward and involves a number of factors.
- Need to take into account the workload, data model, schema design, logical organization of data records in storage devices.
- Extends on the approach of the referenced paper to cover the drawback it had.

Modules and Execution Flow

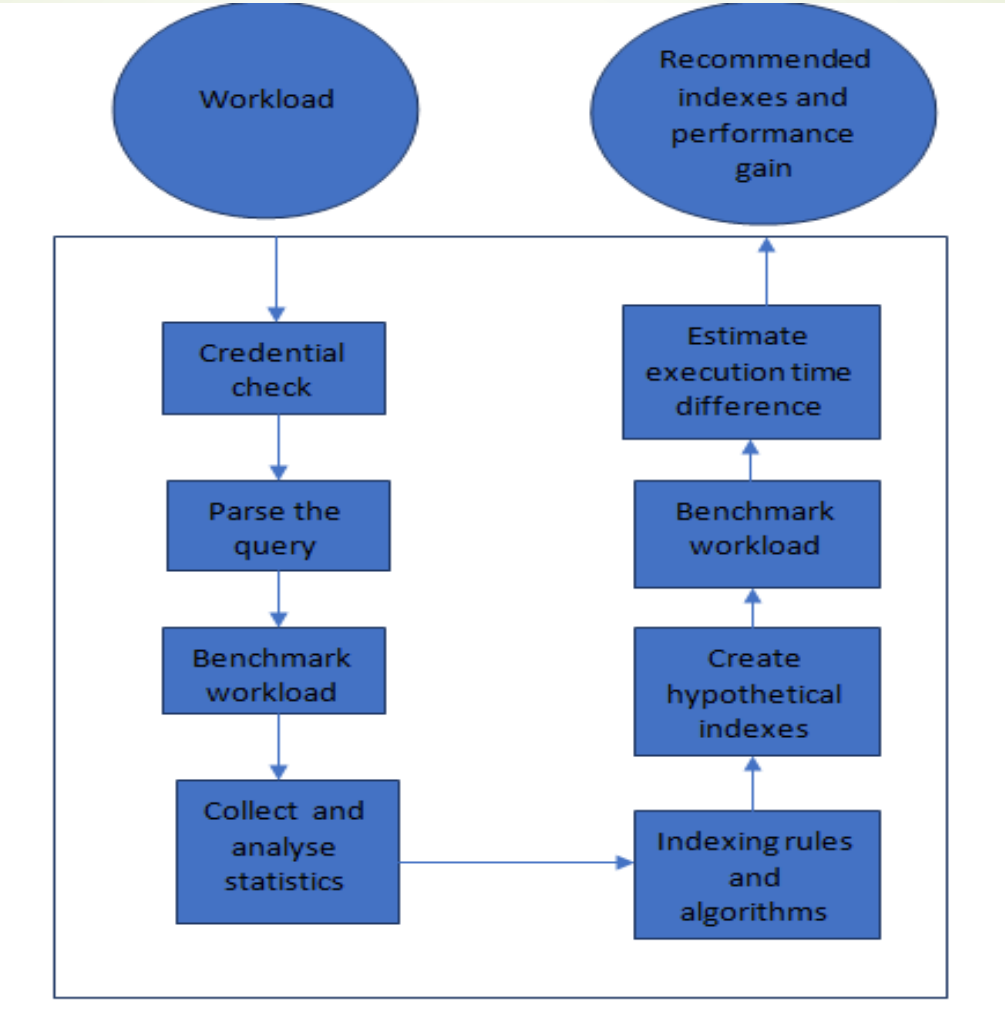

- **Workload:** The input to this project will be the workload – a set of queries that must be triggered against the database chosen.

- **Check admin privilege:** We check if the user has the permission to trigger those queries in the workload and we also validate the credentials of the user with the database

- **Parse the queries:** We parse the queries to develop the statistical data for each of the attributes in the queries.

- **Execute the workload and benchmark:** We execute the queries in the workload against the database with the system generated existing indexes and capture the statistics for each query.

- **Analyze the Workload:** We traverse through the workload and collect the various statistics like the selectivity, reduction factors, predicates, type of data, repetition of the columns, keys, etc., for the attributes and relations involved.

- **Create Hypothetical Indexes:** Based on the statistics, we rank columns of each relation to find out the columns for which index creation would seem useful. We create these indexes for the relation.

- **Bench mark the workload, again:** To quantify the success of the intuitive indexes created, we benchmark the workload again and record I/O statistics.

Index Suggestion Algorithm

- For each table:
  - For each column:
    - For each Clause:
      - If the column occurs for that specific clause we assign it a specific score based on the clause score and number of occurrences.
      - If column occurs in Where Clause:
        - Assign operator score to column depending on the operator it is associated with and number of occurrences.
        - Assign column a hash score or binary score based on association with operators.
  - Find top scoring column and binary and hash score of that column:
    - If binary score is higher → Create binary index
    - If hash index is higher → Check if difference in score between binary and hash score is more than tolerance threshold.
    - If yes → Create hash index

- Composite index:
  - If Binary Index:
    - Check all other columns of table whose binary score is higher and based on a tolerance threshold we add the column to our composite binary index suggestion.
  - If Hash Index:
    - Check all other columns of table whose hash score is higher and based on a tolerance threshold we add the column to our composite hash index suggestion.

•Clause Score
- •Select Clause – 1
- •Having Clause – 2
- •GroupBy, OrderBy – 3
- •Join – 4
- •Where – 5

•Operator Score
- •Equal – 5
- •Greater than Less than – 4
- •Like, Between – 3
- •Not Equal – 2
- •Unknown - 1

Future Works

- Graphical interface which will allow database administrators to check hypothetical performance of database
- Monitor the database workload and create or remove indexes autonomously, effectively providing self-tuning capabilities to the database system.
- Make the tool available online.
- Support for different Database Management Systems.